

Assignment 1

Jeetesh Mangwani, 10314
Harshit Maheshwari, 10290
Vinit Kataria, 10807

POSIX Message Queues

mqd_t mq_open(const char *name, int oflag);

- creates a new message queue or opens an existing queue, returning a message queue descriptor for use in later calls
- returns message queue descriptor of the queue opened if successful
- mqd_t is typically an int
- name is null-terminated string of length less than 256
- oflag is bitwise OR of O_RDONLY (read only mode), O_WRONLY (write only mode), O_RDWR (read-write mode), O_CREAT (create queue if doesn't exist), O_EXCL (if passed with O_CREAT, create queue exclusively if doesn't exist), O_NONBLOCK (open in a non-blocking mode), where the first three are pairwise mutually exclusive

int mq_send(mqd_t mqdes, const char *msg_ptr, size_t msg_len, unsigned msg_prio);

- write a message to a queue
- returns 0 if successful; if fails, returns -1 and sets errno to the code of the error that occurred
- mqdes is the message queue descriptor obtained from mq_open()
- msg_ptr is the pointer to the buffer containing the intended message
- msg_len is the length of the message including NULL terminator
- msg_prio is the message priority
- Doesn't block if the queue was opened with O_NONBLOCK

ssize_t mq_receive(mqd_t mqdes, char *msg_ptr, size_t msg_len, unsigned *msg_prio);

- ssize_t stands for signed size
- if successful, returns the no of bytes excluding the null; if fails, returns -1 and sets errno to the code of the error that occurred
- mqdes is same as in the above function
- msg_ptr is the pointer to the buffer in which to write the message in the front of the queue
- msg_len is the length of the msg_ptr buffer
- if msg_prio is not NULL, the contents of *msg_prio are changed to message priority

int mq_close(mqd_t mqdes);

- closes a message queue that was previously opened
- if successful, returns 0; if fails, returns -1 and sets the errno to code of the error that occurred
- mqdes is the descriptor of the message queue to be closed

int mq_unlink(const char *name);

- removes the message queue named "name" and deallocates it
- if successful, returns 0; if fails, returns -1 and sets the errno to the code of the error that occurred
- name is the pointer to the place containing the name of the queue to be unlinked

Our implementation

- `src/threads/mqueue.c`, `src/threads/mqueue.h`
 - Find the details in the comments too
 - Contains message queue specific data structures, functions, error codes and flags
 - All `sys_mq_*`() functions are same as `mq_*`() functions as described above
 - All the functions are same as POSIX interface except for the following functions:
 - `void mqueue_init(void)`
 - initializes the global static struct list `list_of_message_queues`
 - initializes the message priority comparison function
 - called by `main()` in `src/threads/init.c`
 - `struct mqueue *issue_mqueue(const char* name)`
 - allocates memory for & initializes a fresh message queue; returns a pointer to `mq_open()`, it's caller
 - `bool (*prio_less_than)(const struct list_elem *a, const struct list_elem *b, void* aux)`
 - serves as a pointer to the message-priority comparison function
 - returns true if the message corresponding to a has priority less than that of b
 - aux may be unused
 - The data structures are
 - `enum mq_open_mode`: enumerates the 6 flags for `mq_open`; they are powers of 2 so as to retrieve individual flags from the bitwise or of passed flags
 - `enum mqueue_error_codes`: enumerates the error codes for all the 5 functions described in POSIX section
 - `static char error_names[64][128]`: contains detailed description for each error code
 - `static struct list_of_message_queues`: contains all the message queues of the system
 - `typedef int mqd_t`: serves as identifier for queues
 - `static unsigned int no_of_message_queues`: contains the number of allocated message queues in memory
 - `struct mqueue`:
 - `mqd_t id`: the descriptor of the queue
 - `char name[256]`: the name of the queue
 - `struct list list_of_messages`: contains the messages of this queue
 - `struct list_elem elem`: part of list implementation
 - `int no_of_messages`: no of messages in the queue
 - `struct list read_only_threads, write_only_threads, read_write_threads`: list of threads who have opened the queue in read, write and read-write modes
 - `struct list send_wait_list, rcv_wait_list`: contains entries for threads waiting on this queue to send/receive messages due to queue being full/empty
 - `struct message`
 - `char* text`: contains a pointer to the buffer containing the message
 - `size_t length`: message length
 - `unsigned int priority`: message priority
 - `struct list_elem`: part of list implementation, see `src/lib/kernel/list.c`
 - `struct wait_list_entry`: serves as member of the struct `send_wait_list` and `rcv_wait_list`
 - `struct thread* thread_ptr`: contains the pointer to the thread waiting to send/receive
 - `struct list_elem elem`: part of list implementation
 - `struct tid_holder`: serves as member for the struct `read_only_threads, write_only_threads, read_write_threads` list of the struct `mqueue`
 - `tid_t tid`: thread id of the thread associated
 - `bool blocking`: true if opened in blocking mode
 - `struct list_elem elem`: part of list implementation
 - `src/userprog/syscall.c`
 - added code to initialize `syscall_map[]` and `syscall_narg[]` arrays in `syscall_init()` [which is called from `main` in `src/threads/init.c`]
 - added code to `syscall_handler()` so as to transfer control to the function registered for input `syscall` code
 - `src/lib/user/syscall.c`
 - defined a new macro `syscall4(NUMBER, ARG0, ARG1, ARG2, ARG3)` on similar lines to those defined already
 - `src/lib/syscall-nr.h`
 - added entries `SYS_MQOPEN`, `SYS_MQSEND`, `SYS_MQRECV`, `SYS_MQCLOSE`, `SYS_MQUNLINK` to the enum

Working

Following happens when a message queue function `mq_*`() is called in a thread/process

- `mq_*`() are defined in `src/lib/user/syscall.c`
- `mq_w()` returns `syscallN(SYS_MQW, N arguments...)`, where N is the number of arguments and W is one of {OPEN, SEND, RECEIVE, CLOSE, UNLINK} and w is one of {open, send, receive, close, unlink}

- `syscallN(SYS_MQW, N arguments)` is a macro defined in `src/lib/usr/syscall.c`, where `N` is 0 to 4 [We have implemented `syscall4(NUMBER, ARG0, ARG1, ARG2, ARG3)` as `mq_send()` and `mq_receive()` require four arguments]
- `syscallN()` raises an interrupt with the code `0x30`. The interrupt handler transfers the control to `syscall_handler()` defined in `src/userprog/syscall.c`
- `syscall_handler()` uses the `syscall_map[]` and `syscall_narg[]` arrays initialised by `syscall_init()` [which is called in `main()` in `src/threads/init.c`]
- `syscall_handler(struct intr_frame *f)` transfers the control to the function registered for the number present at `(uint64_t *)(&f->vec_no) + 3` [this was told by TA Vikar Kushwah]
- `typedef int (*handler)(uint32_t, ...):` defines the handler function pointers
- `static handler syscall_map[30]:` contains the map from code to handler
- `static int syscall_narg[30]:` contains the no of arguments for function corr to code
- `syscall_handler()` retrieves the particular function pointer from `syscall_map[]` array
- In our case, the handlers are `sys_mq_*` defined in `src/threads/mqueue.c`
- in case of any error, `sys_mq_*` sets the `current_thread()->error` to one of the codes defined in `mqueue.h` [we have added an `int error` to struct `thread` in `src/threads/threads.h`] and returns `-1` or `0` as the definition dictates

Test files

- We have made two test files : `src/tests/threads/message-passing-block.c` and `src/tests/threads/message-passing-nonblock.c` for testing blocking and non-blocking calls respectively
- Find the details in the comments
- To run, type the following commands in `src/threads/build`
 - `make clean`
 - `make`
 - `pintos run message-passing-block`
 - `pintos run message-passing-nonblock`